

Thesis Abstract: Improving Multi-Object Navigation in an LLM-Based Embodied AI Framework

Mohammadkazem Rajabi

Introduction

Large language models have shown remarkable abilities in planning, reasoning, and solving complex tasks by generating code. In recent work such as ViperGPT[1] and VisProg[2], LLMs are used to generate complex compositional programs for language-vision-tasks like object detection or visual question answering, allowing them to solve complex visual tasks without any training. These systems show significant results in vision-language challenges just by using a few in-context examples to create complex programs by combining different modules without training. Figure 1 illustrates how VisProg leverages an LLM to generate multi-step visual programs from natural language prompts.

Embodied AI[3] focuses on building agents that can move, see, and reason within 3D environments essentially learning through interaction with their environments rather than relying solely on static data. Embodied AI has seen significant progress in recent years, largely thanks to the development of fast, photorealistic simulators [4, 5]. These tools allow researchers to run large numbers of parallel simulations in realistic indoor environments, making it possible to test and evaluate agents at scale something that was much harder in traditional robotics setups.

Related Work

Alongside these advances, a wide range of embodied tasks and benchmarks has been introduced. These include classic navigation goals like PointGoal[6], ObjectGoal [7](finding a target object), and Instance-Image Goal (locating an object from an image), as

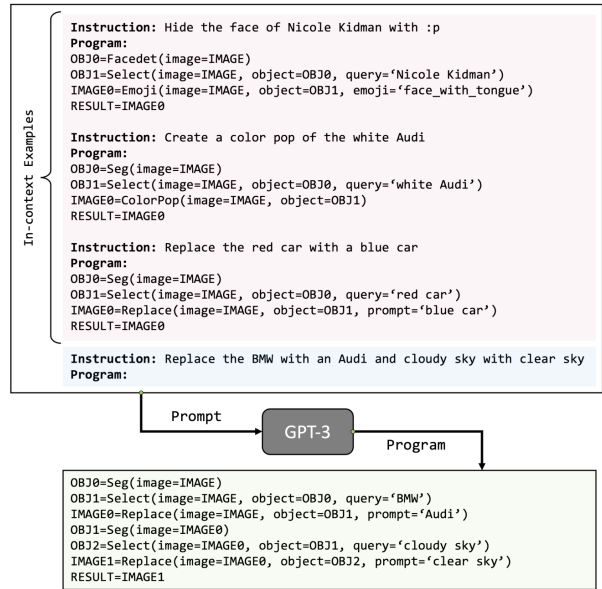


Figure 1: Example from VisProg examples [1].

well as more complex tasks Embodied Question Answering (EQA)[8], where an AI agent must understand an environment well enough to answer open-ended natural language questions.

To solve these tasks, researchers have explored end-to-end models[9], modular systems[10], and, approaches driven by large language models[11]. However, many of these methods still rely on task-specific architectures. End-to-end models often require massive amounts of training to perform well, and adapting them to new tasks can be slow and costly.

TANGO Framework

TANGO[12] is a framework in embodied AI that leverages LLMs’ capabilities to generate programs; this work seeks to extend this ability to embodied agents, enabling them to observe and control their actions in zero-shot navigation rather than learning task-specific policies. The complete set of TANGO modules, including navigation, vision-language, and functional subroutines, is illustrated in Figure 2.




Figure 2: Overview of TANGO modules.

The agent receives a natural language prompt (e.g., ”Is the shower curtain open?”), The LLM generates a structured program composed of high-level commands like detect, navigate-to, or answer and integrates it into the Habitat-Sim framework[13].

As shown in Figure 3, the LLM is guided by a fixed set of in-context examples, which act as a prompt template. It depicts how to break a complex task into smaller steps.

The generated program is passed to the TANGO’s program interpreter, which maps the generated code to each module for execution; the LLM only generates high-level planning, and the program interpreter executes it. For example, the detect() is mapped to the whole detect module, executes the full code, and returns the output. In this way, we can detect failures in specific modules and enhance them. This design allows TANGO to generalize across a wide range of embodied challenges.

Improving Object Detection

The target of this project is to enhance a bottleneck in TANGO; for some tasks, the accuracy and




Figure 3: Overview of program generation and execution in TANGO.

success rate are low due to failures in the detection part, specifically in MultiON[6] tasks. The baseline open-set detector, OWL-ViT2[14], showed poor performance in these cases, prompting a focused effort to improve the detection pipeline.

We tackle this problem by fine-tuning several state-of-the-art detection models, including YOLOv8 [15] and DETR [16] variants on our synthetic dataset for MultiON task using Habitat-Sim and HM3D scenes. We got an improvement by using our best-performing model, YOLOv8-X, in the detection module.

When evaluated on the MultiON 3-ON task, the upgraded detector improved the agent’s success rate by 28% compared to the OWL-ViT2 baseline. It also outperformed the MOPA framework [17] by 14.6% in success and exceeded the best end-to-end model, ProjNeuralMap [6], by 26%. These results depict the point of integrating fine-tuned detectors in the detection module.

Despite this limitation, our TANGO-based approach achieves significantly excellent results, demonstrating its effectiveness in generalizable and zero-shot navigation in MultiON without retraining.

Additionally, ablation studies on TANGO’s memory mechanism are conducted to understand its power in agent performance by creating a map and remembering the visited areas instead of searching for objects. The results confirmed improvements across

both 3-ON and 2-ON settings when memory was used in combination with the upgraded detector. We summarize the performance of various YOLO variants and baselines on the MultiON 3-ON task in Table 1.

Model	Success (%)	Progress
PredSem [17]	38	53
No-Map [17]	10	24
ObjRecogMap [17]	22	40
ProjNeuralMap [17]	27	46
YOLO8n With Memory	48	67
YOLO8m With Memory	49	69
YOLO8l With Memory	47	67.3
YOLO8x With Memory	52.6 (+28%)	70 (+27%)
Baseline (Owl-ViT2) With Memory	24	43
YOLO8x Without Memory	49	68.8
Baseline (Owl-ViT2) Without Memory	19	39

Table 1: Performance Comparison on MultiON 3-ON Task

In summary, this work demonstrates that using improved object detectors and a memory mechanism significantly boosts the performance of the MultiON task. By leveraging large language models for program generation, the system achieves robust zero-shot generalization in embodied AI.

References

- [1] Amey Gupta et al. “VisProg: A Neuro-Symbolic Program Generator for Visual Reasoning”. In: *arXiv preprint arXiv:2211.11559* (2023).
- [2] Karan Desai et al. “ViperGPT: Visual Perception using Language Constructs”. In: *arXiv preprint arXiv:2303.08128* (2023).
- [3] Matt Deitke et al. “Retrospectives on the Embodied AI Workshop”. In: *arXiv preprint arXiv:2210.06849* (2022). URL: <https://arxiv.org/abs/2210.06849>.
- [4] Xavier Puig et al. “Habitat 3.0: A Co-Habitat for Humans, Avatars and Robots”. In: *arXiv preprint arXiv:2310.13724* (2023). Available at <https://arxiv.org/abs/2310.13724>. URL: <https://arxiv.org/abs/2310.13724>.
- [5] Santhosh K Ramakrishnan and et al. “HM3D: A large-scale 3D dataset for indoor scene understanding”. In: *arXiv preprint arXiv:2103.11778* (2021).
- [6] Saim Wani et al. “MultiON: Benchmarking Semantic Map Memory using Multi-Object Navigation”. In: *arXiv preprint arXiv:2012.03912* (2020). Available at <https://arxiv.org/abs/2012.03912>.
- [7] Dhruv Batra et al. “ObjectNav Revisited: On Evaluation of Embodied Agents Navigating to Objects”. In: *arXiv preprint arXiv:2006.13171* (2020). Available at <https://arxiv.org/abs/2006.13171>.
- [8] Arjun Majumdar et al. “OpenEQA: Embodied Question Answering in the Era of Foundation Models”. In: *Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024.
- [9] Tommaso Campari et al. “Exploiting Scene-specific Features for Object Goal Navigation”. In: *arXiv preprint arXiv:2008.09403* (2020).
- [10] Jacob Krantz et al. “Navigating to Objects Specified by Images”. In: *arXiv preprint arXiv:2304.01192* (2023).
- [11] Vishnu Sashank Dorbala, James F. Mullen Jr, and Dinesh Manocha. “Can an Embodied Agent Find Your ‘Cat-shaped Mug’? LLM-Guided Exploration for Zero-Shot Object Navigation”. In: *arXiv preprint arXiv:2303.03480* (2023). Also published in IEEE Robotics and Automation Letters, Jan 2023.
- [12] Filippo Ziliotto et al. “TANGO: Training-free Embodied AI Agents for Open-world Tasks”. In: *arXiv preprint arXiv:2412.10402* (2024).
- [13] Manolis Savva* et al. “Habitat: A Platform for Embodied AI Research”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2019.

- [14] Neil Houlsby Matthias Minderer Alexey Gritsenko. “Scaling Open-Vocabulary Object Detection”. In: *NeurIPS* (2023).
- [15] Joseph Redmon et al. “You Only Look Once: Unified, Real-Time Object Detection”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition (CVPR)*. 2016, pp. 779–788. URL: <https://arxiv.org/abs/1506.02640>.
- [16] Nicolas Carion et al. “End-to-End Object Detection with Transformers”. In: *ECCV*. 2020.
- [17] Sonia Raychaudhuri et al. “MOPA: Modular Object Navigation With PointGoal Agents”. In: *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. Jan. 2024, pp. 5763–5773.